

MOVIE RECOMMENDER

Business Understanding

Users often struggle to choose what to watch. Personalized movie recommendations can help cut through the noise, offering tailored suggestions that enhance user satisfaction, boost engagement, and encourage longterm use. This project aims to build a recommendation system to simulate how such platforms deliver personalized experiences.

Data Preparation

In [55]:

```
# Installing necessary libraries  
!pip install numpy==1.23.5  
!pip install scikit-surprise  
!pip install wordcloud
```

```
Requirement already satisfied: numpy==1.23.5 in /usr/local/lib/python3.11/dist-packages (1.23.5)  
Requirement already satisfied: scikit-surprise in /usr/local/lib/python3.11/dist-packages (1.1.4)  
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/dist-packages (from scikit-surprise) (1.4.2)  
Requirement already satisfied: numpy>=1.19.5 in /usr/local/lib/python3.11/dist-packages (from scikit-surprise) (1.23.5)  
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.11/dist-packages (from scikit-surprise) (1.14.1)  
Requirement already satisfied: wordcloud in /usr/local/lib/python3.11/dist-packages (1.9.4)  
Requirement already satisfied: numpy>=1.6.1 in /usr/local/lib/python3.11/dist-packages (from wordcloud) (1.23.5)  
Requirement already satisfied: pillow in /usr/local/lib/python3.11/dist-packages (from wordcloud) (11.1.0)  
Requirement already satisfied: matplotlib in /usr/local/lib/python3.11/dist-packages (from wordcloud) (3.10.0)  
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->wordcloud) (1.3.2)  
Requirement already satisfied: cyclor>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib->wordcloud) (0.12.1)  
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib->wordcloud) (4.57.0)  
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->wordcloud) (1.4.8)  
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib->wordcloud) (24.2)  
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->wordcloud) (3.2.3)  
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.11/dist-packages (from matplotlib->wordcloud) (2.8.2)  
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.7->matplotlib->wordcloud) (1.17.0)
```

In []:

```
# Loading the relevant libraries  
  
import pandas as pd  
import seaborn as sns  
import matplotlib.pyplot as plt  
import numpy as np  
from scipy.stats import skew, kurtosis
```

```
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LinearRegression, Lasso
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import mean_squared_error
from scipy.stats import linregress
from sklearn.impute import SimpleImputer
from sklearn.pipeline import make_pipeline
from sklearn.metrics import r2_score
from sklearn.model_selection import train_test_split
from surprise import Dataset, Reader, SVD
from sklearn.metrics.pairwise import cosine_similarity
from sklearn.preprocessing import MultiLabelBinarizer
```

```
#NLP Libraries
from sklearn.feature_extraction.text import CountVectorizer
import nltk
from nltk.corpus import stopwords
nltk.download('stopwords')

#Wordcloud to visualize most frequent terms in the tags
from wordcloud import WordCloud
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
```

In [57]:

```
# Loading the data
# Load specific CSV files inside the folder
movies = pd.read_csv('/content/movies.csv')
ratings = pd.read_csv('/content/ratings.csv')
tags = pd.read_csv('/content/tags.csv')
links = pd.read_csv('/content/links.csv')
```

Data Understanding

In [58]:

```
movies.head()
```

Out[58]:

movieId		title	genres
0	1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy
1	2	Jumanji (1995)	Adventure Children Fantasy
2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama Romance
4	5	Father of the Bride Part II (1995)	Comedy

In [59]:

```
ratings.head()
```

Out[59]:

	userId	movieId	rating	timestamp
0	1	1	4.0	964982703
1	1	3	4.0	964981247
2	1	6	4.0	964982224
3	1	47	5.0	964983815

```
4  userId 1  movieId 50  rating 5.0  timestamp 964982931
```

In [60]:

```
links.head()
```

Out[60]:

	movieid	imdbid	tmdbid
0	1	114709	862.0
1	2	113497	8844.0
2	3	113228	15602.0
3	4	114885	31357.0
4	5	113041	11862.0

In [61]:

```
tags.head()
```

Out[61]:

	userid	movieId	tag	timestamp
0	2	60756	funny	1445714994
1	2	60756	Highly quotable	1445714996
2	2	60756	will ferrell	1445714992
3	2	89774	Boxing story	1445715207
4	2	89774	MMA	1445715200

In [62]:

```
movies.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9742 entries, 0 to 9741
Data columns (total 3 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   movieId    9742 non-null   int64
 1   title      9742 non-null   object
 2   genres     9742 non-null   object
dtypes: int64(1), object(2)
memory usage: 228.5+ KB
```

In [63]:

```
ratings.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100836 entries, 0 to 100835
Data columns (total 4 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   userId     100836 non-null  int64
 1   movieId    100836 non-null  int64
 2   rating     100836 non-null  float64
 3   timestamp  100836 non-null  int64
dtypes: float64(1), int64(3)
memory usage: 3.1 MB
```

In [64]:

```
tags.info()
```

```
colnames(lmmodel$aov.frame$DataFrame)\
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3683 entries, 0 to 3682
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   userId      3683 non-null   int64
1   movieId     3683 non-null   int64
2   tag         3683 non-null   object
3   timestamp   3683 non-null   int64
dtypes: int64(3), object(1)
memory usage: 115.2+ KB
```

In [65]:

```
links.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9742 entries, 0 to 9741
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   movieId     9742 non-null   int64
1   imdbId      9742 non-null   int64
2   tmdbId      9734 non-null   float64
dtypes: float64(1), int64(2)
memory usage: 228.5 KB
```

Data Cleaning

Checking for missing values

In [66]:

```
movies.isnull().sum()
```

Out[66]:

	0
movieId	0
title	0
genres	0

dtype: int64

In [67]:

```
ratings.isnull().sum()
```

Out[67]:

	0
userId	0
movieId	0
rating	0
timestamp	0

dtype: int64

In [68]:

```
tags.isnull().sum()
```

Out[68]:

	0
userId	0
movieId	0
tag	0
timestamp	0

dtype: int64

In [69]:

```
links.isnull().sum()
```

Out[69]:

	0
movieId	0
imdbId	0
tmdbId	8

dtype: int64

In [70]:

```
#Dropping the missing values.

links.dropna(subset=['tmdbId'], inplace=True)

links.isnull().sum()
```

Out[70]:

	0
movieId	0
imdbId	0
tmdbId	0

dtype: int64

Data Preprocessing

In [71]:

```
#Converting the timestamp column into datetime format

ratings['timestamp'] = pd.to_datetime(ratings['timestamp'], unit='s')
ratings.head()
```

Out[71]:

	userId	movieId	rating	timestamp
0	1	1	4.0	2000-07-30 18:45:03
1	1	3	4.0	2000-07-30 18:20:47
2	1	6	4.0	2000-07-30 18:37:04
3	1	47	5.0	2000-07-30 19:03:35
4	1	50	5.0	2000-07-30 18:48:51

In [72]:

```
#Converting the timestamp column into datetime format

tags['timestamp'] = pd.to_datetime(tags['timestamp'], unit='s')
tags.head()
```

Out[72]:

	userId	movieId	tag	timestamp
0	2	60756	funny	2015-10-24 19:29:54
1	2	60756	Highly quotable	2015-10-24 19:29:56
2	2	60756	will ferrell	2015-10-24 19:29:52
3	2	89774	Boxing story	2015-10-24 19:33:27
4	2	89774	MMA	2015-10-24 19:33:20

In [73]:

```
#Extract the year from the title column

movies['year'] = movies['title'].str.extract(r'\((\d{4})\)')
movies['year'] = movies['year'].fillna(0).astype(int)
```

In [74]:

```
movies.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9742 entries, 0 to 9741
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype  
---  -
0   movieId     9742 non-null   int64  
1   title       9742 non-null   object  
2   genres      9742 non-null   object  
3   year        9742 non-null   int64  
dtypes: int64(2), object(2)
memory usage: 304.6+ KB
```

In [75]:

```
movies.head()
```

Out[75]:

	movieId	title	genres	year
0	1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy	1995
1	2	Jumanji (1995)	Adventure Children Fantasy	1995
2	3	Grumpier Old Men (1995)	Comedy Romance	1995
3	4	Waiting to Exhale (1995)	Comedy Drama Romance	1995
4	5	Father of the Bride Part II (1995)	Comedy	1995

In [76]:

```
# Merge ratings and movies

df = pd.merge(ratings, movies, on='movieId')

df.head()
```

Out[76]:

userId	movieId	rating	timestamp	title		genres	year
0	1	1	4.0	2000-07-30 18:45:03	Toy Story (1995)	Adventure Animation Children Comedy Fantasy	1995
1	1	3	4.0	2000-07-30 18:20:47	Grumpier Old Men (1995)	Comedy Romance	1995
2	1	6	4.0	2000-07-30 18:37:04	Heat (1995)	Action Crime Thriller	1995
3	1	47	5.0	2000-07-30 19:03:35	Seven (a.k.a. Se7en) (1995)	Mystery Thriller	1995
4	1	50	5.0	2000-07-30 18:48:51	Usual Suspects, The (1995)	Crime Mystery Thriller	1995

Merged ratings and movies because we will need genres, titles, and ratings together to do feature engineering and modeling.

Extracted Year from the title of the movie because older movies might be rated differently. Movie age is useful for modeling.

Converted the timestamp column into date time format to make more sense for this analysis.

EDA

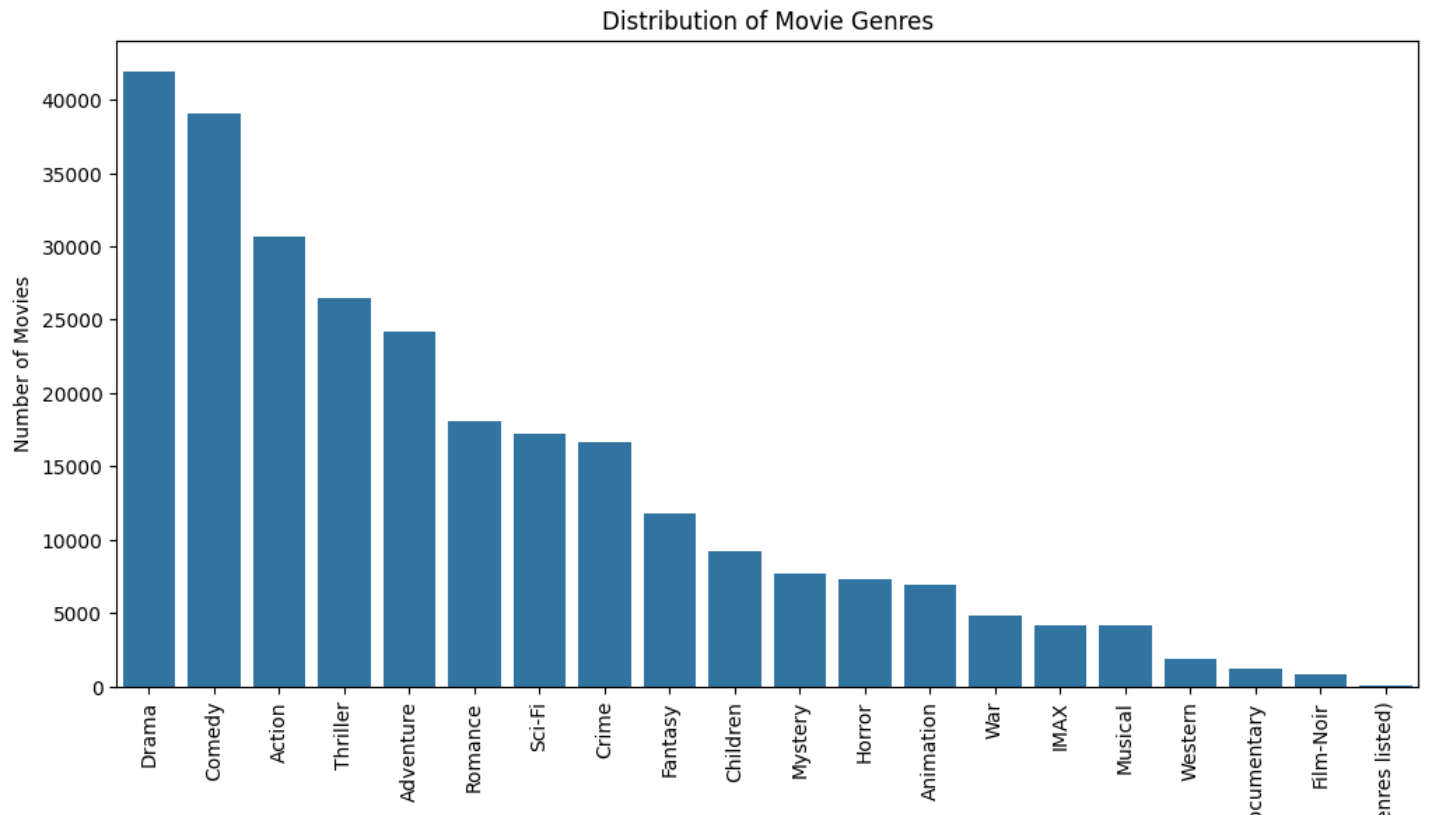
In [78]:

```
# Analyze Top genre distribution

genre_counts = df['genres'].str.split('|').explode().value_counts()

plt.figure(figsize=(12, 6))
sns.barplot(x=genre_counts.index, y=genre_counts.values)
plt.xticks(rotation=90)
plt.title('Distribution of Movie Genres')
plt.xlabel('Genre')
plt.ylabel('Number of Movies')
plt.show()

print("Most Common Genres:")
print(genre_counts.head(10))
```



Most Common Genres:

```
genres
Drama      41928
Comedy     39053
Action     30635
Thriller   26452
Adventure  24161
Romance    18124
Sci-Fi     17243
Crime      16681
Fantasy    11834
Children   9208
Name: count, dtype: int64
```


In [79]:

```
# Movie popularity

movie_popularity = df['movieId'].value_counts()
sns.histplot(movie_popularity, bins=50, kde=False)
plt.title('Ratings per Movie')
```

Out[79]:

Text(0.5, 1.0, 'Ratings per Movie')



In [80]:

```
# Number of Unique Users and Movies

print("Number of unique users:", df['userId'].nunique())
print("Number of unique movies:", df['movieId'].nunique())
```

Number of unique users: 610
Number of unique movies: 9724

In [81]:

```
# Ratings Distribution
```



```
print("Rating distribution:")
df['rating'].value_counts().sort_index()
```

Rating distribution:


Out[81]:

	count
rating	
0.5	1370
1.0	2811
1.5	1791
2.0	7551
2.5	5550
3.0	20047
3.5	13136
4.0	26818
4.5	8551
5.0	13211

dtype: int64

In [82]:

```
df['rating'].value_counts().sort_index().plot(kind='bar')
plt.title('Rating Distribution')
plt.xlabel('Rating')
plt.ylabel('Count')
plt.show()
```



In [83]:

```
# Ratings per user
ratings_per_user = df.groupby('userId')['rating'].count()
print("Average ratings per user:", ratings_per_user.mean())
```

```
# Ratings per movie
ratings_per_movie = df.groupby('movieId')['rating'].count()
print("Average ratings per movie:", ratings_per_movie.mean())
```

Average ratings per user: 165.30491803278687
Average ratings per movie: 10.369806663924312

In [84]:

```
print("Top users by number of ratings:")
print(ratings_per_user.sort_values(ascending=False).head())

print("Top movies by number of ratings:")
top_movies = ratings_per_movie.sort_values(ascending=False).head()
print(top_movies)
```

Top users by number of ratings:
userId
414 2698
599 2478
474 2108
448 1864
274 1346
Name: rating, dtype: int64
Top movies by number of ratings:
movieId
356 329
318 317
296 307
593 279
2571 278
Name: rating, dtype: int64

In [85]:

```
most Rated = movies.groupby('title').size().sort_values(ascending=False).head(10)
print(most Rated)
```

```
title
Confessions of a Dangerous Mind (2002)      2
Emma (1996)                                  2
Eros (2004)                                  2
War of the Worlds (2005)                     2
Saturn 3 (1980)                              2
Partisan (2015)                              1
Parenthood (1989)                            1
Paris Is Burning (1990)                      1
Paris, I Love You (Paris, je t'aime) (2006)  1
Paris, Texas (1984)                          1
dtype: int64
```

Feature Engineering

One Hot-Encode Genres

In []:

```
# Splitting genres

df['genre_list'] = df['genres'].apply(lambda x: x.split('|'))

mlb = MultiLabelBinarizer()
genre_encoded = pd.DataFrame(mlb.fit_transform(df['genre_list']), columns=mlb.classes_)
df = pd.concat([df, genre_encoded], axis=1)
```

In [87]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100836 entries, 0 to 100835
Data columns (total 48 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   userId                               100836 non-null  int64
1   movieId                             100836 non-null  int64
2   rating                              100836 non-null  float64
3   timestamp                           100836 non-null  datetime64[ns]
4   title                               100836 non-null  object
5   genres                              100836 non-null  object
6   year                                100836 non-null  int64
7   genre_list                          100836 non-null  object
8   (no genres listed)                 100836 non-null  int64
9   Action                             100836 non-null  int64
10  Adventure                           100836 non-null  int64
11  Animation                           100836 non-null  int64
12  Children                            100836 non-null  int64
13  Comedy                              100836 non-null  int64
14  Crime                               100836 non-null  int64
15  Documentary                         100836 non-null  int64
16  Drama                               100836 non-null  int64
17  Fantasy                             100836 non-null  int64
18  Film-Noir                          100836 non-null  int64
19  Horror                              100836 non-null  int64
20  IMAX                                100836 non-null  int64
21  Musical                             100836 non-null  int64
22  Mystery                             100836 non-null  int64
23  Romance                             100836 non-null  int64
24  Sci-Fi                              100836 non-null  int64
25  Thriller                            100836 non-null  int64
26  War                                 100836 non-null  int64
27  Western                             100836 non-null  int64
28  (no genres listed)                 100836 non-null  int64
29  Action                             100836 non-null  int64
30  Adventure                           100836 non-null  int64
31  Animation                           100836 non-null  int64
32  Children                            100836 non-null  int64
33  Comedy                              100836 non-null  int64
34  Crime                               100836 non-null  int64
35  Documentary                         100836 non-null  int64
36  Drama                               100836 non-null  int64
37  Fantasy                             100836 non-null  int64
38  Film-Noir                          100836 non-null  int64
39  Horror                              100836 non-null  int64
40  IMAX                                100836 non-null  int64
41  Musical                             100836 non-null  int64
42  Mystery                             100836 non-null  int64
43  Romance                             100836 non-null  int64
44  Sci-Fi                              100836 non-null  int64
45  Thriller                            100836 non-null  int64
46  War                                 100836 non-null  int64
47  Western                             100836 non-null  int64
dtypes: datetime64[ns](1), float64(1), int64(43), object(3)
memory usage: 36.9+ MB
```

In [88]:

```
df.head()
```

Out[88]:

userId	movieId	rating	timestamp	title	genres	year	genre_list	(no genres listed)
0	1	1	2000-07-30 18:45:03	Toy Story (1995)	Adventure Animation Children Comedy Fantasy	1995	[Adventure, Animation, Children, Comedy, Fantasy]	0
			2000-07-	Grumpier				

1	1	3	4.0	2000-07-30	Old Men (1993)	Comedy/Romance	1995	[Comedy, Romance]	(no genres listed)	Ac
userId	movieId	rating	timestamp	timeStamp	title	genres	year	genre_list	genres	
2	1	6	4.0	2000-07-30 18:37:04	Heat (1995)	Action/Crime/Thriller	1995	[Action, Crime, Thriller]	0	
3	1	47	5.0	2000-07-30 19:03:35	Seven (a.k.a. Se7en) (1995)	Mystery/Thriller	1995	[Mystery, Thriller]	0	
4	1	50	5.0	2000-07-30 18:48:51	Usual Suspects, The (1995)	Crime/Mystery/Thriller	1995	[Crime, Mystery, Thriller]	0	

5 rows x 48 columns



In [89]:

```
df.drop(columns=['(no genres listed)'], inplace = True)
```

In [90]:

```
df.drop(columns=['genre_list'], inplace=True)
```

In [91]:

```
# Filter Rarely Rated Movies / Users (Cold Start Filtering)
# For modeling, remove users/movies with too few ratings

def filter_rare_interactions(df, user_threshold, movie_threshold):
    """
    Filters out users and movies with a low number of interactions.

    Args:
        df: DataFrame containing user-movie interactions.
        user_threshold: Minimum number of movies rated by a user.
        movie_threshold: Minimum number of ratings for a movie.

    Returns:
        Filtered DataFrame.
    """

    # Count user and movie interactions
    user_counts = df['userId'].value_counts()
    movie_counts = df['movieId'].value_counts()

    # Identify users and movies to keep
    active_users = user_counts[user_counts >= user_threshold].index
    popular_movies = movie_counts[movie_counts >= movie_threshold].index

    # Filter the DataFrame
    filtered_df = df[
        (df['userId'].isin(active_users)) & (df['movieId'].isin(popular_movies))
    ]

    return filtered_df

user_threshold = 5
movie_threshold = 5
filtered_ratings = filter_rare_interactions(ratings, user_threshold, movie_threshold)

print(f"Original DataFrame shape: {ratings.shape}")
print(f"Filtered DataFrame shape: {filtered_ratings.shape}")
```

Original DataFrame shape: (100836, 4)
Filtered DataFrame shape: (90274, 4)

Merging Ratings with Movie Info to help associate each rating with a movie title and genre.

In [92]:

```
df_ratings_movies = pd.merge(ratings, movies, on='movieId', how='left')
print(df_ratings_movies.head())
```

	userId	movieId	rating	timestamp	title \
0	1	1	4.0	2000-07-30 18:45:03	Toy Story (1995)
1	1	3	4.0	2000-07-30 18:20:47	Grumpier Old Men (1995)
2	1	6	4.0	2000-07-30 18:37:04	Heat (1995)
3	1	47	5.0	2000-07-30 19:03:35	Seven (a.k.a. Se7en) (1995)
4	1	50	5.0	2000-07-30 18:48:51	Usual Suspects, The (1995)

	genres	year
0	Adventure Animation Children Comedy Fantasy	1995
1	Comedy Romance	1995
2	Action Crime Thriller	1995
3	Mystery Thriller	1995
4	Crime Mystery Thriller	1995

In [93]:

```
# Calculate movie popularity
movie_popularity = filtered_ratings.groupby('movieId')['userId'].count().reset_index()
movie_popularity.columns = ['movieId', 'rating_count']


# Define popularity thresholds (adjust as needed)
popularity_threshold = movie_popularity['rating_count'].median()
movie_popularity['popularity_level'] = movie_popularity['rating_count'].apply(
    lambda x: 'Popular' if x >= popularity_threshold else 'Obscure'
)


# Merge popularity information with movie details
movie_popularity = pd.merge(movie_popularity, movies[['movieId', 'title']], on='movieId',
                             how='left')

# Visualization: Movie Popularity Distribution
plt.figure(figsize=(10, 6))
sns.histplot(movie_popularity['rating_count'], bins=30, kde=True)
plt.title('Distribution of Movie Rating Counts')
plt.xlabel('Number of Ratings')
plt.ylabel('Number of Movies')
plt.show()

#Further analysis and visualization
popular_movies = movie_popularity[movie_popularity['popularity_level'] == 'Popular']
obscure_movies = movie_popularity[movie_popularity['popularity_level'] == 'Obscure']
```

Distribution of Movie Rating Counts





Create the user-item matrix to represents user preferences explicitly.

In [94]:

```
# Create the user-movie rating matrix
ratings_matrix = filtered_ratings.pivot_table(index='userId', columns='movieId', values='rating')

# Fill missing values with 0 (indicating no rating)
ratings_matrix = ratings_matrix.fillna(0)

print(ratings_matrix.head())
```

movieId	1	2	3	4	5	6	7	8	\
userId									
1	4.0	0.0	4.0	0.0	0.0	4.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
5	4.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	

movieId	9	10	...	176371	177593	177765	179401	179819	180031	\
userId										
1	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	
2	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	
4	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	
5	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	

movieId	180985	183897	187593	187595
userId				
1	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0
5	0.0	0.0	0.0	0.0

[5 rows x 3650 columns]

Looking at how engaged users are


In [95]:

```
# Calculate user engagement metrics
user_engagement = filtered_ratings.groupby('userId')['movieId'].count().reset_index()
user_engagement.columns = ['userId', 'interaction_count']

# Define engagement categories
engagement_threshold = user_engagement['interaction_count'].median()
user_engagement['engagement_level'] = user_engagement['interaction_count'].apply(lambda x : 'Frequent' if x >= engagement_threshold else 'Occasional')

# Visualization: Distribution of Interaction counts
plt.figure(figsize=(10, 6))
sns.histplot(user_engagement['interaction_count'], bins=30, kde=True)
plt.title('Distribution of User Interaction Counts')
plt.xlabel('Number of Interactions')
plt.ylabel('Number of Users')
plt.show()
```

Distribution of User Interaction Counts



Identifying popular vs obscure movies so as to help with understanding cold-start issues

In [96]:

```
# Calculate movie popularity
movie_popularity = filtered_ratings.groupby('movieId')['userId'].count().reset_index()
movie_popularity.columns = ['movieId', 'rating_count']

# Define popularity thresholds (adjust as needed)
popularity_threshold = movie_popularity['rating_count'].median()
movie_popularity['popularity_level'] = movie_popularity['rating_count'].apply(
    lambda x: 'Popular' if x >= popularity_threshold else 'Obscure'
)

# Merge popularity information with movie details
movie_popularity = pd.merge(movie_popularity, movies[['movieId', 'title']], on='movieId',
                             how='left')

# Visualization: Movie Popularity Distribution
plt.figure(figsize=(10, 6))
sns.histplot(movie_popularity['rating_count'], bins=30, kde=True)
plt.title('Distribution of Movie Rating Counts')
plt.xlabel('Number of Ratings')
plt.ylabel('Number of Movies')
plt.show()

# Further analysis and visualization
popular_movies = movie_popularity[movie_popularity['popularity_level'] == 'Popular']
obscure_movies = movie_popularity[movie_popularity['popularity_level'] == 'Obscure']

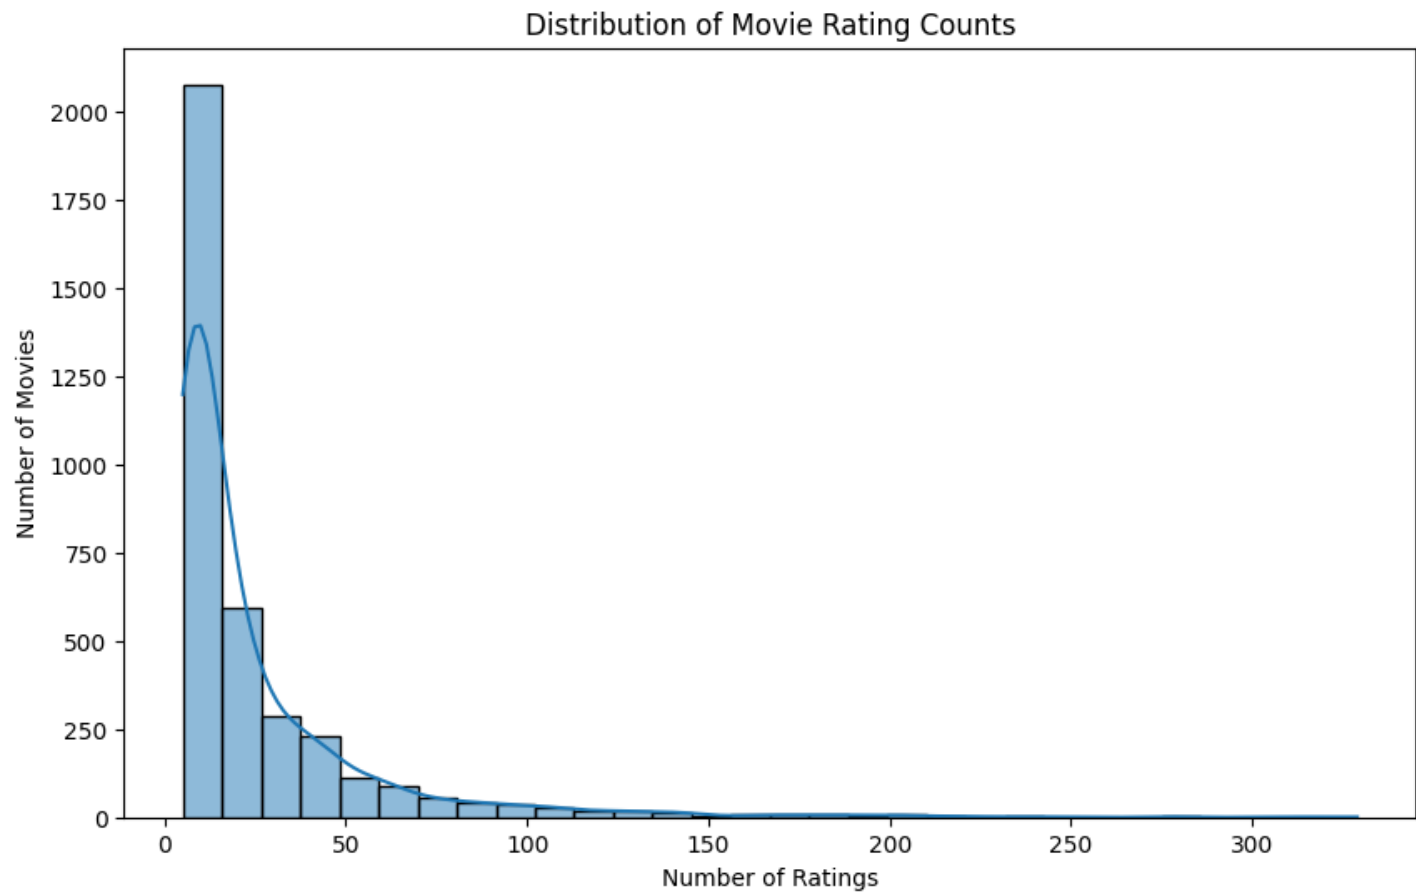
print("\nPopular Movies (Examples):")
print(popular_movies.head())

print("\nObscure Movies (Examples):")
print(obscure_movies.head())

# Average rating of popular vs. obscure movies visualization
average_ratings = filtered_ratings.groupby('movieId')['rating'].mean().reset_index()
movie_popularity = pd.merge(movie_popularity, average_ratings, on='movieId', how='left')

plt.figure(figsize=(8, 6))
```

```
sns.boxplot(x='popularity_level', y='rating', data=movie_popularity)
plt.title('Average Rating of Popular vs. Obscure Movies')
plt.xlabel('Popularity Level')
plt.ylabel('Average Rating')
plt.show()
```





Popular Movies (Examples):

	movieId	rating_count	popularity_level	title
0	1	215	Popular	Toy Story (1995)
1	2	110	Popular	Jumanji (1995)
2	3	52	Popular	Grumpier Old Men (1995)
4	5	49	Popular	Father of the Bride Part II (1995)
5	6	102	Popular	Heat (1995)

Obscure Movies (Examples):

	movieId	rating_count	popularity_level	title
3	4	7	Obscure	Waiting to Exhale (1995)
7	8	8	Obscure	Tom and Huck (1995)
12	13	8	Obscure	Balto (1995)
26	27	9	Obscure	Now and Then (1995)
27	28	11	Obscure	Persuasion (1995)






Heatmap of Correlation Between Movie Ratings

In [97]:

```
popular_movies = movie_popularity.nlargest(20, 'rating_count')['movieId']
ratings_subset = ratings_matrix[popular_movies]

# Calculate the correlation matrix
correlation_matrix = ratings_subset.corr()

plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Heatmap of Correlation Between Movie Ratings (Popular Movies)')
plt.show()
```



Modeling

Collaborative Filtering - Matrix Factorization (SVD)

In [119]:

```
# Create a Surprise Dataset
reader = Reader(rating_scale=(0.5, 5))
data = Dataset.load_from_df(filtered_ratings[['userId', 'movieId', 'rating']], reader)

# Train the SVD model
trainset = data.build_full_trainset()
svd_model = SVD(n_factors=100, n_epochs=20, lr_all=0.005, reg_all=0.02)
svd_model.fit(trainset)

# Function to get top N recommendations for a user
def get_top_n_recommendations(user_id, n=5):
    # Get all movies the user has not rated
    user Rated movies = set(filtered_ratings[filtered_ratings['userId'] == user_id]['movieId'])
    all_movies = set(filtered_ratings['movieId'].unique())
    unrated_movies = list(all_movies - user Rated movies)

    # Predict ratings for unrated movies
    predictions = []
    for movie_id in unrated_movies:
        prediction = svd_model.predict(user_id, movie_id)
        predictions.append((movie_id, prediction.est))

    # Sort predictions by estimated rating and get top N
    predictions.sort(key=lambda x: x[1], reverse=True)
    top_n = predictions[:n]

    # Return movie titles instead of IDs
    top_n_movies = []
    for movie_id, rating in top_n:
        movie_title = movies[movies['movieId']==movie_id]['title'].values[0]
        top_n_movies.append((movie_title, rating))

    return top_n_movies

user_id = 1
recommendations = get_top_n_recommendations(user_id)
print(f"Top 5 movie recommendations for user {user_id}:")
for movie, rating in recommendations:
    print(f"- {movie} (Predicted rating: {rating:.2f})")
```

Top 5 movie recommendations for user 1:

- City of Lost Children, The (Cité des enfants perdus, La) (1995) (Predicted rating: 5.00)
- Hoop Dreams (1994) (Predicted rating: 5.00)
- Shawshank Redemption, The (1994) (Predicted rating: 5.00)
- Wallace & Gromit: A Close Shave (1995) (Predicted rating: 5.00)
- Trainspotting (1996) (Predicted rating: 5.00)

Content-Based Filtering (Using Movie Metadata)

In [106]:

```
# Create a TF-IDF matrix from the movie genres
tfidf = CountVectorizer()
tfidf_matrix = tfidf.fit_transform(movies['genres'])

# Compute the cosine similarity between movies based on the TF-IDF matrix
cosine_sim = cosine_similarity(tfidf_matrix, tfidf_matrix)
```

```

# Function to get movie recommendations based on genre similarity
def content_based_recommendations(movie_title, cosine_sim=cosine_sim):
    # Get the index of the movie
    idx = movies.index[movies['title'] == movie_title].tolist()[0]

    # Get the pairwise similarity scores
    sim_scores = list(enumerate(cosine_sim[idx]))

    # Sort the movies based on similarity scores
    sim_scores = sorted(sim_scores, key=lambda x: x[1], reverse=True)

    # Get the top 5 most similar movies
    sim_scores = sim_scores[1:6]

    # Get the movie indices
    movie_indices = [i[0] for i in sim_scores]

    # Return the top 5 recommended movies
    return movies['title'].iloc[movie_indices]

# using Toy Story (1995) as an example but you can replace it with the movie title you want
recommendations = content_based_recommendations('Toy Story (1995)')
recommendations

```

Out[106]:

	title
1706	Antz (1998)
2355	Toy Story 2 (1999)
2809	Adventures of Rocky and Bullwinkle, The (2000)
3000	Emperor's New Groove, The (2000)
3568	Monsters, Inc. (2001)

dtype: object

This code is implementing a content-based movie recommendation system using genre similarity so as to recommend movies that are similar in genre to a given movie, hence returning a list of 5 movies that have the most similar genres to the movie. we decided to use "Toy Story (1995)" as an example.

Hybrid models

In [120]:

```

def hybrid_recommendations(user_id, n_recommendations=5):
    """
    Generate hybrid recommendations by combining collaborative filtering and content-based filtering.

    Args:
        user_id: The ID of the user for whom to generate recommendations.
        n_recommendations: The number of recommendations to generate.

    Returns:
        A list of movie titles representing the recommendations.
    """
    # Get collaborative filtering recommendations
    collab_recs = collaborative_recommendations(user_id, n_recommendations * 2)

    # Get content-based recommendations for each of the collab movies
    content_based_candidates = []
    for movie_id, _ in collab_recs:
        content_based_candidates.extend(content_based_recommendations(movie_id))

    # Combine and rank

```

```

hybrid_scores = {}
for movie_id, score in collab_recs:
    hybrid_scores[movie_id] = score

for movie_id in content_based_candidates:
    if movie_id not in hybrid_scores:
        # Estimate rating using SVD for unseen movies in content recommendations
        prediction = svd.predict(user_id, movie_id)
        hybrid_scores[movie_id] = prediction.est

sorted_hybrid = sorted(hybrid_scores.items(), key=lambda x: x[1], reverse=True)

# Get movie titles instead of IDs:
recommended_movie_titles = []
for movie_id, _ in sorted_hybrid[:n_recommendations]:
    title = movies[movies['movieId'] == movie_id]['title'].iloc[0]
    recommended_movie_titles.append(title)

return recommended_movie_titles # Return movie titles

# using user_ID = 1 as an example but you replace with the user ID you want recommendations for
user_id = 1
recommended_movies = hybrid_recommendations(user_id)
recommended_movies

```

Out[120]:

```

['Shawshank Redemption, The (1994)',
 'Lion King, The (1994)',
 'In the Name of the Father (1993)',
 'Blade Runner (1982)',
 'Wallace & Gromit: The Best of Aardman Animation (1996)']

```

MODEL EVALUATION

Model evaluation using RMSE

In [121]:

```

from sklearn.metrics import mean_squared_error
import math

def rmse_evaluation(predictions):
    """
    Calculates the Root Mean Squared Error (RMSE) for a list of predictions.

    Args:
        predictions: A list of Prediction objects from Surprise library.

    Returns:
        The RMSE value.
    """
    true_ratings = [pred.r_ui for pred in predictions]
    estimated_ratings = [pred.est for pred in predictions]
    rmse = math.sqrt(mean_squared_error(true_ratings, estimated_ratings))
    return rmse

from surprise.model_selection import train_test_split

# Assuming 'reader' is the reader object defined previously
reader = Reader(rating_scale=(0.5, 5))
data = Dataset.load_from_df(filtered_ratings[['userId', 'movieId', 'rating']], reader=reader)

trainset, testset = train_test_split(data, test_size=0.25, random_state=42)

# Train an SVD model

```

```
algo = SVD()
algo.fit(trainset)

# Make predictions on the test set
predictions = algo.test(testset)

# Calculate and print the RMSE
rmse = rmse_evaluation(predictions)
print(f"RMSE: {rmse}")
```

RMSE: 0.8577033705647412

Model evaluation using MAE

In [123]:

```
from sklearn.metrics import mean_absolute_error

def mae_evaluation(predictions):
    """
    Calculates the Mean Absolute Error (MAE) for a list of predictions.

    Args:
        predictions: A list of Prediction objects from Surprise library.

    Returns:
        The MAE value.
    """
    true_ratings = [pred.r_ui for pred in predictions]
    estimated_ratings = [pred.est for pred in predictions]
    mae = mean_absolute_error(true_ratings, estimated_ratings)
    return mae

# Assuming 'predictions' is the list of predictions from the SVD model
mae = mae_evaluation(predictions)
print(f"MAE: {mae}")
```

MAE: 0.6586869449672015

Conclusions and recommendations

In []: